

PROJECT OF AAPP

K-mer counting



Diego Gaboardi and Giorgio Giardini
Politecnico di Milano

PROBLEM

Counting the number of occurrences of every k-mer in a long string

K-mer → substring of length k

USAGE

genome assembly

alphabet = {A,C,G,T}

Ref.	ATGCTGATGCTAGCGATATGCCCTAAAATCGATGCTAGCTGACTGATCGATCGACTGTCA
Read 1	CCTAAAATCG
Read 2	CCCTAAAATC
Read 3	GCCCTAAAAT
Read 4	TGCCCTAAAA
Read 5	ATGCCCTAAA
Read 6	TATGCCCTAA
Read 7	ATATGCCCTA
Read 8	GATATGCCCT
Read 9	CGATATGCC
Read 10	GCGATATGCC

statistics



knowledge



STANDARD IMPLEMENTATION

Hash table

Key: substring (k-mers)

Value: counter

Key	CCTA	ACGT	TTGC	ATCC	AATG	GAAC	TCCG	AGTC	CCAG	TGAA
Value	567	234	456	567	234	567	123	673	578	198

SPATIAL DIMENSION (in bit)

$$A^k * \log_2(b)$$

Where A is the alphabet size

k the substrings length


b the maximum value of each counter

TEMPORAL DIMENSION

$$|S| * \text{table_access_time}$$

Where S is the input string

ON MAP REDUCE



Every node process
a substring

Straightforward implementation -> map operation emits a series of couples $\langle \text{k-mer}, 1 \rangle$ for each node -> large overhead

Efficient implementation -> map operation includes a counter and emits a series of couples $\langle \text{k-mer}, \text{count} \rangle$



How to obtain efficiently?

Jellyfish parallel algorithm on
each node

JELLYFISH

➤ GOALS:

- ❑ Fast and multithreaded
- ❑ Memory efficient



➤ SOLUTION:


- ❑ Lock free Hash-table with CAS assembly instruction
- ❑ Merging intermediate hash tables
- ❑ Reduced memory usage of hash entry
- ❑ Space-efficient encoding of keys



PARALLELISM: SIMD

- Single Instruction
 - Same code with a loop cycle
 - Each iteration of the loop is executed by one of the threads in the team
- Multiple Data
 - Partitioning of the string between the threads
 - Each thread process a contiguous sequence of characters

! **Border issue:** overlapping at border



OUR IMPLEMENTATION

Language: C++ & open mp

- Support shared memory multiprocessing
- Easy to use simply adding some commands to standard C++ code

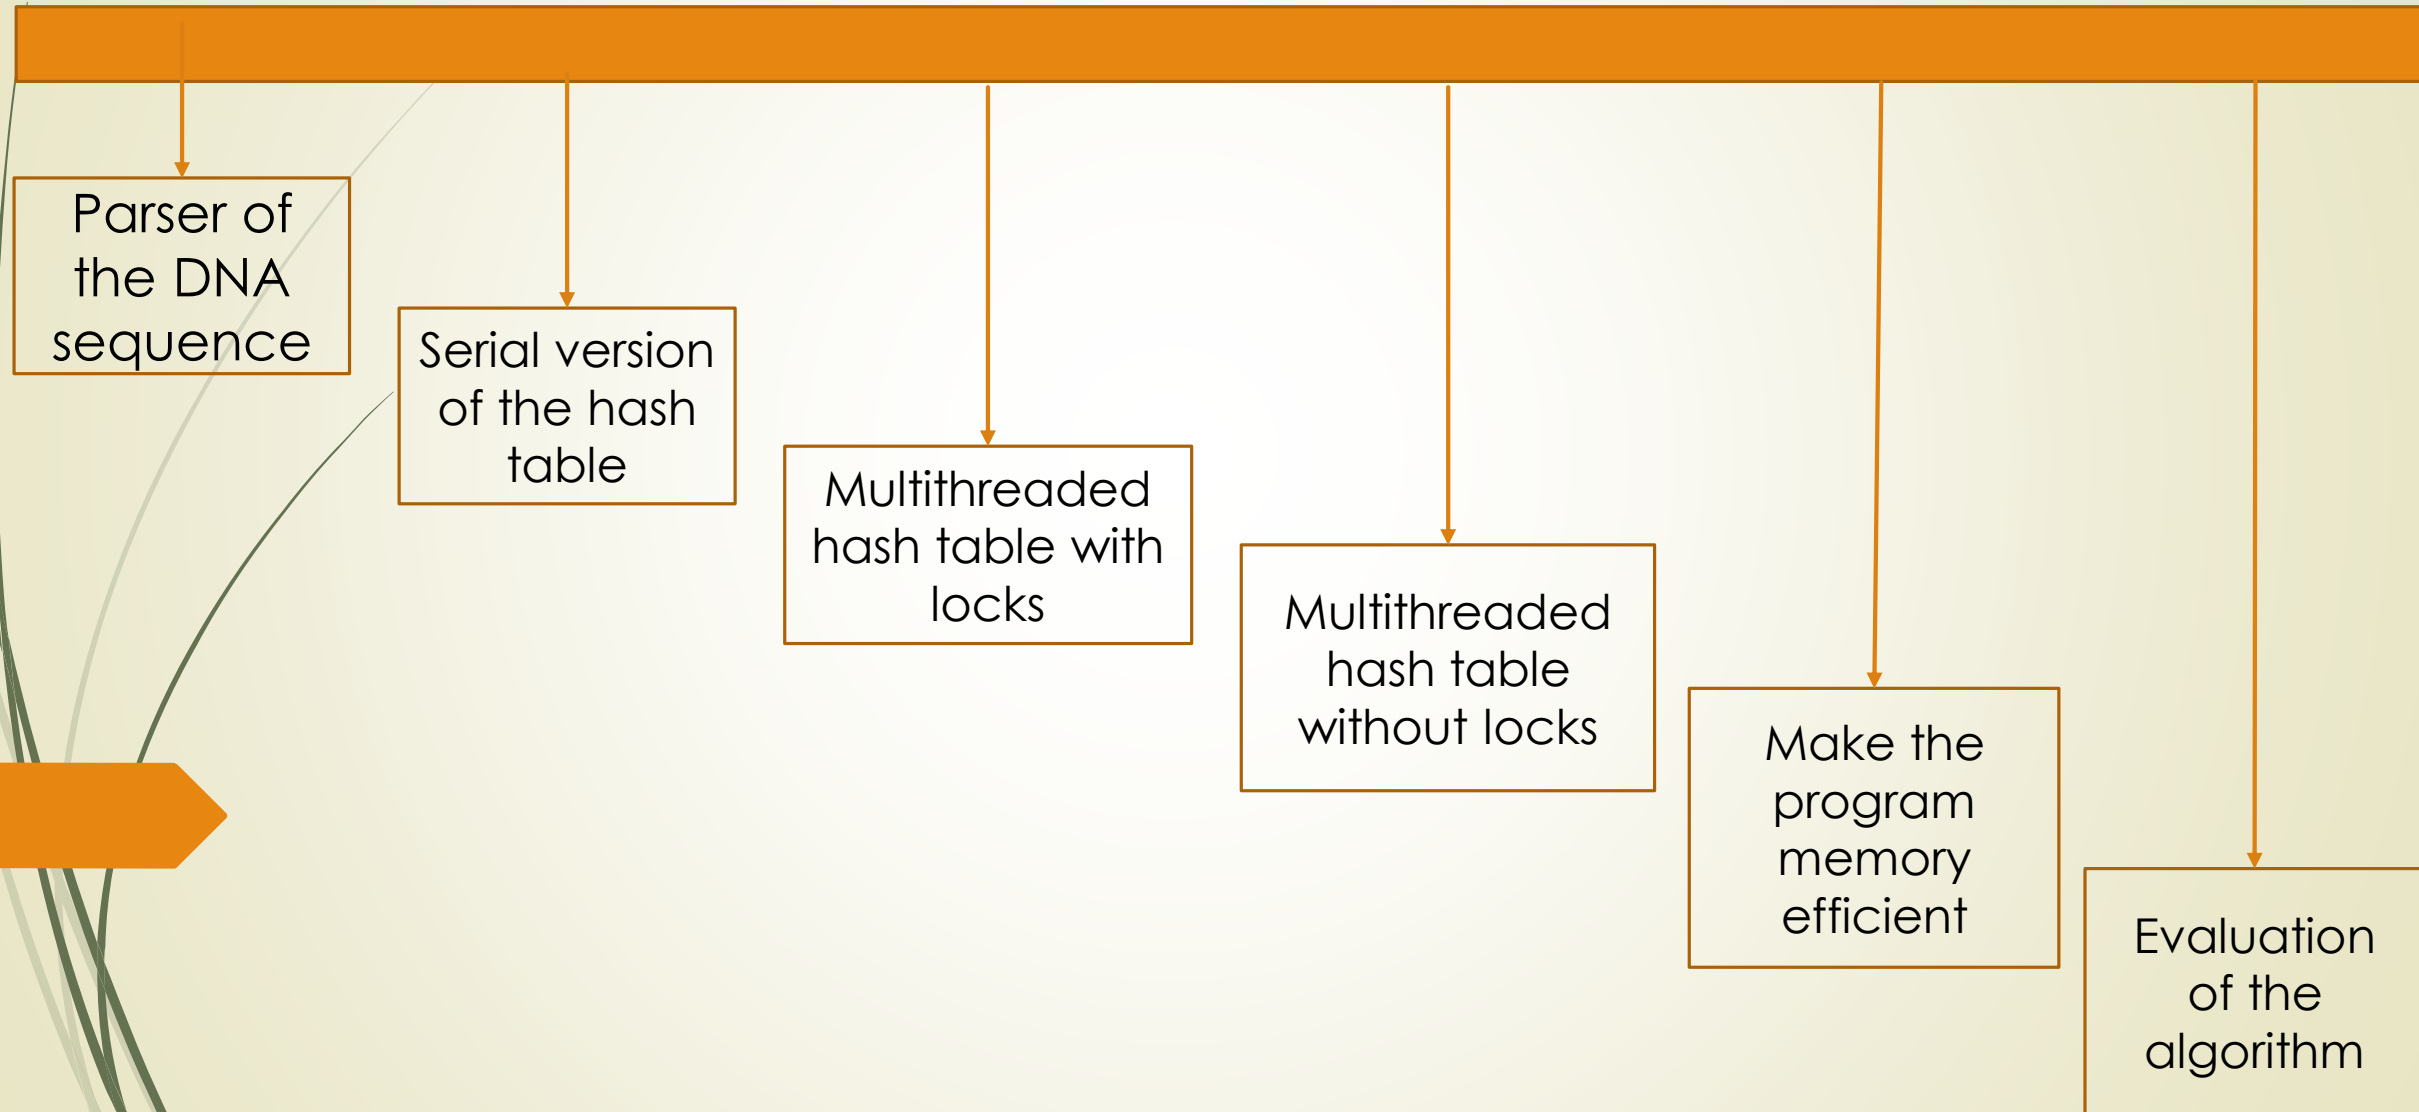
GIT and GIT HUB

- For working in group

C++




WORK PLAN





OUR IMPLEMENTATION



SERIAL IMPLEMENTATION

➤ Class model

- HashTable
- HashEntry
- Structures

```
/*a single entry of the hash table*/
class HashEntry
{
    private:
        std::vector<Nucleotide> key;
        int count;
    public:
        HashEntry();
        bool isEmpty();
        int getC();
        std::vector<Nucleotide> getK();
        void setC(int);
        void setK(std::vector<Nucleotide>);
        std::string toString();
};
```

```
/*class used to implement the hash table in which are counted the k-mers*/
class HashTable
{
    private:
        int k, L;
        std::vector<HashEntry> table;
        MatrixXf matrix, inverse;
        int f(std::vector<Nucleotide>);
        int reprobe(int);
    public:
        HashTable(int, int);
        void incrementValue(std::vector<Nucleotide>);
        std::string toString();
};
```

```
/*A Nucleotide: A, C, G or T*/
class Nucleotide
{
    private:
        std::bitset<2> nucleotide;
    public:
        Nucleotide(char);
        string toString();
        bool equal(Nucleotide);
        int getBit(int);
};
```

SERIAL IMPLEMENTATION

➤ Program option

```
po::options_description desc;

desc.add_options()
    ("h, help", "Shows description of the options")
    ("l, lenght", po::value<int>(&l)->default_value(10), "Set the lenght of the hash table; default value 10.")
    ("t, thread", po::value<int>(&thread)->default_value(24), "Set the number of threads")
    ("f, file", po::value<std::string>(&file_name)->default_value("../dna_sequences/DNA_prova.txt"),
     "Set the name of the file in which there is the dna sequence; default value ../dna_sequences/DNA_prova.txt.")
    ("k, k_lenght", po::value<int>(&k_lenght)->default_value(4), "Set the lenght of k; default value 4.");

po::variables_map vm;
po::store(po::parse_command_line(argc, argv, desc), vm);
po::notify(vm);
omp_set_num_threads(thread);
```

Program options let user set the following parameters at execution time:

- l – hash table lenght
- t – number of thread
- f – source file
- k – k-mer lenght

LOCK MULTITHREAD IMPLEMENTATION

OUR CHOICE: use of a critical call of increment function, in an openmp parallel cycle

```
/*cycle in which the k-mer are added to the hash table and so counted*/
#pragma omp parallel for
for(int i=0; i<dna_sequence.size()-k_lenght; i++)
{
    std::vector<Nucleotide> k_mer;
    for(int j=i; j<i+k_lenght; j++)
    {
        k_mer.push_back(dna_sequence[j]);
    }
    #pragma omp critical
    hashTable.incrementValue(k_mer);
}
```



OPTIMIZATIONS





LOCK-FREE HASH TABLE

CAS instruction:

- Reads a memory location
- Compare the read value with the second parameter of the instruction
- If they are equal write the memory with the 3rd parameter
- Return the previously held value

Possibility to detect simultaneous access to shared resources

USED TO:

- Finds the location in the hash table
- Increments the value associated with the key



IMPLEMENTATION

OUR CHOICE: use the GCC function `__sync_bool_compare_and_swap()`

- Find the location in the hash table

```
int i=0;
int pos;
bool done = false;

do{
    pos = (hash + HashTable::reprobe(i)) % m;
    i++;
    done = __sync_bool_compare_and_swap(&table_hash[pos],0,key_value);
    if(key_value == table_hash[pos])
        done = true;
    if(table_count[pos] >= MAX_COUNT-1)
        done = false;
}
while(!done && i < MAX_REPROBE );
```


- Increments the value associated with the key

```
done = false;
int oldCount = table_count[pos];

do{
    oldCount = table_count[pos];
    done = __sync_bool_compare_and_swap(&table_count[pos],oldCount,oldCount+1);
}while(!done);
```



MERGING INTERMEDIATE HASH TABLE

- Once computed the hash table is written to disk in order
 - In this way it is possible to:
 - Query quickly using binary search
 - Merge two hash table
 - If there isn't enough memory for the entire computation:
 - Intermediate results are saved to disk
 - The hash table is cleared and we begin counting afresh
 - At the end intermediate results are merged
- 



IMPLEMENTATION

OUR CHOICE: use an openmp critical section in order to avoid writing conflicts
order the table efficiently using quicksort algorithm

```
bool is_max_reprobe = false;
if(i>=MAX_REPROBE)
#pragma omp critical
//block in which the content of the hash table is saved on the disk and the hash table is flushed
{
    stop = true; //bool used to block the other threads

    //the file is created
    std::string file_name = "result";
    file_name.append(std::to_string(num));
    std::ofstream outfile(file_name);

    //the hash table is ordered and then written on the file
    this->order(0, this->m);
    outfile << this->toString();
    outfile.close();

    //the hash table is flushed
    num++;
    this->flush();
    stop = false;
    is_max_reprobe = true;

    //the key count is now updated
    this->incrementValue(key);
}

if(is_max_reprobe)
    return;
```

```
/*method that order the hashTable according to quicksort algorithm*/
void HashTable::order(int low, int high)
{
    if (low < high)
    {
        //pi is partitioning index, arr[p] is now at right place
        int pi = partition_hash(low, high);

        // Separately sort elements before partition and after partition
        order(low, pi - 1);
        order(pi + 1, high);
    }
}
```

REDUCED MEMORY USAGE

- ▶ Using a value field large enough for the most repeated k-mer: **INEFFICIENT**
- ▶ Most of the k-mers appear:
 - 1 → sequencing error
 - C → sequencing coverage
- ▶ **IDEA:**
 - Small value field
 - Allow key to have more than one entry



IMPLEMENTATION

- Small value field

Definition of count attribute as an int (32bit)

```
std::vector<int> table_count;  
#define MAX_COUNT 4294967285 //max value of count
```

- Allow keys to have more than one entry

```
do{  
    pos = (hash + HashTable::reprobe(i)) % m;  
    i++;  
    done = __sync_bool_compare_and_swap(&table_hash[pos],0,key_value);  
    if(key_value == table_hash[pos])  
        done = true;  
    if(table_count[pos] >= MAX_COUNT-1)  
        done = false;  
}  
while(!done && i < MAX_REPROBE );
```

KEY ENCODING

$M = 2^l$  length of hash-table

$$\text{pos}(m, i) = (\text{hash}(m) + \text{reprobe}(i)) \bmod M$$

The position of the key give as information about the l -lower bits

IDEA: store only 2^{k-l} higher bits and the reprobe count

In many application we can obtain a space per key that is independent of the length of the l -mers and of the input string



IMPLEMENTATION

OUR CHOICE: do not implement, not fully compatible with other optimizations

Why?

The optimization is based on the fact that the key is encoded in the smallest possible field.

In order to take advantage from this optimization we must use a key value field that is the minimum large to encode a k-mer.

The dimension of this field is variable and depends on k.

The CAS function is an assembly primitive. It does not work with non-primitive data type.

Primitive data types have all fixed length, minimum 16bit for short int.

The only way to implement this optimization is to use CAS with short int, but this does not improve so much the performance of the algorithm so we do not include this optimization in our final implementation

TIME EVALUATION

- ▶ 3 Version of the algorithm:

- Serial version
- Multithreaded hash-table with locks
- Multithreaded hash-table without locks

- ▶ We want to:

- Show the speed up with respect to the number of threads
- How much time it takes with respect to the lenght of the sequence

- ▶ We did each execution of the programs three times and then calculate the average




DATASET

- **DNA randomly generated**
- **gbbct155.seq**: complete genome of Escherichia Coli
- **gbbct367.seq**: complete genome of Lactococcus lactis
- **gbgss116.seq**: Chlorocebus aethiops genomic clone CH252-491O17, genomic survey sequence
- **hs_alt_CHM1_1.1_chr22.gbk**: Homo sapiens chromosome 22 genomic scaffold




RESULT

- Execution time using different datasets




RESULT

- Execution time compared to dataset length (random dna sequence)



RESULT

- Speedup with respect to dataset length (random dna sequence)



RESULT


- Execution time compared to K

$L = 2 * K$

THREADS = 24

DS = gbg116

K



RESULT


- Execution time compared to number of thread

$L = 2 * K$

$K = 4$


DS = gbg116

Number Of Thread




RESULT

- Memory usage compared to dataset length (random dna sequence)



RESULT

- Memory usage compared to K (all versions have a similar behaviour with respect to k)



CONSIDERATIONS

- ▶ Lock-parallel-version has always very bad performances
- ▶ Optimised version is always better than the other algorithms
- ▶ Lock-parallel-version scales bad with the number of threads
- ▶ Optimised version scales well with the number of threads
- ▶ Execution time grows linearly with the input length
- ▶ Speedup grows with the input length
- ▶ Execution time grows linearly with k
- ▶ Memory usage grows linearly with the input length
- ▶ Memory usage grows exponential with k





REFERENCES

[National Center for Biotechnology Information
<https://www.ncbi.nlm.nih.gov/genbank/release/145/>]

[Random DNA Generator
<http://www.faculty.ucr.edu/mmaduro/random.htm>]

[A fast, lock-free approach for efficient parallel counting of occurrences of k-mers Guillaume Marçais¹, and Carl Kingsford]

[GitHub repository
https://github.com/DiegoGabo/ProgettoAAPP/tree/master/dna_sequences]

[Compare and swap documentation
http://www.cplusplus.com/reference/atomic/atomic/compare_exchange_strong/]